

# Reinforcement Learning

## A Quick Introduction

Nir Ben-Zvi

June 11, 2017

# Outline

## 1 What is Reinforcement Learning About

- Reinforcement Learning Examples
- Formal Introduction

## 2 Approaches to Reinforcement Learning

- Model-based RL
- Markov Decision Processes
- Value-based RL
- Bellman Equations
- Partially Observable Markov Decision Processes
- Policy-based RL

## 3 Solving RL in Practice

- Q-Learning
- Policy Gradient Methods
- Policy Gradients Explained
- REINFORCE
- Algorithm Comparison

## 4 Using Policy Gradients with Neural Networks

- Quick Example

## 5 Wrapping Up

- Further Reading/Acknowledgements

# Outline

## 1 What is Reinforcement Learning About

- Reinforcement Learning Examples
- Formal Introduction

## 2 Approaches to Reinforcement Learning

- Model-based RL
- Markov Decision Processes
- Value-based RL
- Bellman Equations
- Partially Observable Markov Decision Processes
- Policy-based RL

## 3 Solving RL in Practice

- Q-Learning
- Policy Gradient Methods
- Policy Gradients Explained
- REINFORCE
- Algorithm Comparison

## 4 Using Policy Gradients with Neural Networks

- Quick Example

## 5 Wrapping Up

- Further Reading/Acknowledgements

# Why is RL Cool?

- RL has been getting a lot of wins in academy in recent years
- Like many other branches of ML which got the 'deep' treatment
- We'll later see why the two fuse together nicely


# What Makes It Different?

- Informally, RL is about processes (or actions)
- It's not strictly supervised or unsupervised learning
- Instead, it's about learning how to interact with a certain process
- Games are one of the famous examples

# Examples


- Atari Games (DeepMind)
- Pong (From Karpathy's Blogpost)

# Example: Learning Pong



taken from Karpathy's Blogpost

# Example: Learning Pong



taken from Karpathy's Blogpost



# Other Non-Game Uses

- Visual Attention Models (Mnih et Al, now DeepMind)
- StarCraft (DeepMind)
- Robotics
- Neural Turing Machines (DeepMind) and followup works
- Alpha Go and games

# Outline

## 1 What is Reinforcement Learning About

- Reinforcement Learning Examples
- **Formal Introduction**

## 2 Approaches to Reinforcement Learning

- Model-based RL
- Markov Decision Processes
- Value-based RL
- Bellman Equations
- Partially Observable Markov Decision Processes
- Policy-based RL

## 3 Solving RL in Practice

- Q-Learning
- Policy Gradient Methods
- Policy Gradients Explained
- REINFORCE
- Algorithm Comparison

## 4 Using Policy Gradients with Neural Networks

- Quick Example

## 5 Wrapping Up

- Further Reading/Acknowledgements

# Agent Environment Interaction

Reinforcement Learning usually includes an environment and an agent

- For the purpose of this intro lecture assume a single agent
- Environment/agent combos can be a casino/gambler, maze/robot etc.

# Agent Environment Interaction

Reinforcement Learning usually includes an environment and an agent

- For the purpose of this intro lecture assume a single agent
- Environment/agent combos can be a casino/gambler, maze/robot etc.
- RL is about learning how to interact with the environment with accordance to timesteps
- Many approaches exist, and we'll present a few today

# Agent Environment Interaction, cont.

- The agent learns what to do so as to maximize a numerical reward
- The learner is not told which actions to take, but instead must discover which actions yield high rewards from experience
- Rewards are not immediate
- These two characteristics - *trial-and-error search* and *delayed reward* are the two most important distinguishing features of RL

# Agent Environment Interaction, cont.

- The agent learns what to do so as to maximize a numerical reward
- The learner is not told which actions to take, but instead must discover which actions yield high rewards from experience
- Rewards are not immediate
- These two characteristics - *trial-and-error search* and *delayed reward* are the two most important distinguishing features of RL

# Outline

## 1 What is Reinforcement Learning About

- Reinforcement Learning Examples
- Formal Introduction

## 2 Approaches to Reinforcement Learning

- Model-based RL
- Markov Decision Processes
- Value-based RL
- Bellman Equations
- Partially Observable Markov Decision Processes
- Policy-based RL

## 3 Solving RL in Practice

- Q-Learning
- Policy Gradient Methods
- Policy Gradients Explained
- REINFORCE
- Algorithm Comparison

## 4 Using Policy Gradients with Neural Networks

- Quick Example

## 5 Wrapping Up

- Further Reading/Acknowledgements

# Reinforcement Learning Approaches

An RL agent usually includes one or more of the following:

- **Model**: agent's understanding of the true environment
- **Value Function**: how good is each state and/or action
- **Policy**: the agent's behavior function



# Outline

## 1 What is Reinforcement Learning About

- Reinforcement Learning Examples
- Formal Introduction

## 2 Approaches to Reinforcement Learning

- Model-based RL**
- Markov Decision Processes
- Value-based RL
- Bellman Equations
- Partially Observable Markov Decision Processes
- Policy-based RL

## 3 Solving RL in Practice

- Q-Learning
- Policy Gradient Methods
- Policy Gradients Explained
- REINFORCE
- Algorithm Comparison

## 4 Using Policy Gradients with Neural Networks

- Quick Example

## 5 Wrapping Up

- Further Reading/Acknowledgements

# Model-based RL

- Build a model of the environment
- Plan (e.g. by lookahead) using model
- We won't discuss this today

# Outline

## 1 What is Reinforcement Learning About

- Reinforcement Learning Examples
- Formal Introduction

## 2 Approaches to Reinforcement Learning

- Model-based RL
- **Markov Decision Processes**
- Value-based RL
- Bellman Equations
- Partially Observable Markov Decision Processes
- Policy-based RL

## 3 Solving RL in Practice

- Q-Learning
- Policy Gradient Methods
- Policy Gradients Explained
- REINFORCE
- Algorithm Comparison

## 4 Using Policy Gradients with Neural Networks

- Quick Example

## 5 Wrapping Up

- Further Reading/Acknowledgements

# Background; Markov Chains

- A Markov Chain (or Markov Process) is a memoryless random process - i.e. a sequence of random states  $s_1, s_2, \dots$  which abide the *Markov Property*

# Background; Markov Chains

- A Markov Chain (or Markov Process) is a memoryless random process - i.e. a sequence of random states  $s_1, s_2, \dots$  which abide the *Markov Property*

## Markov Property

A state  $s_t$  is *Markov* if and only if:

$$\mathbb{P}[s_{t+1} \mid s_t] = \mathbb{P}[s_{t+1} \mid s_1, \dots, s_t]$$

# Background; Markov Chains

- A Markov Chain (or Markov Process) is a memoryless random process - i.e. a sequence of random states  $s_1, s_2, \dots$  which abide the *Markov Property*

## Markov Property

A state  $s_t$  is *Markov* if and only if:

$$\mathbb{P}[s_{t+1} \mid s_t] = \mathbb{P}[s_{t+1} \mid s_1, \dots, s_t]$$

- In other words; the state captures all relevant information about the past
- When a state is known, previous events are meaningless

# State Transition Matrix

- The Markov State transition matrix defines a Markov Process;

$$\mathcal{P}_{ss'} = \mathbb{P}[s_{t+1} = s' \mid s_t = s]$$

- The State Transition matrix  $\mathcal{P}$  defines transition probabilities from all states  $s$  to all successor states  $s'$ :

$$\mathbb{P} = \begin{bmatrix} \mathcal{P}_{11} & \dots & \mathcal{P}_{1n} \\ \vdots & & \vdots \\ \mathcal{P}_{n1} & \dots & \mathcal{P}_{nn} \end{bmatrix}$$

- ...such that rows and columns sum to one
- Also called *doubly stochastic*

# Markov Process, Revisited

We'll now define a Markov Process more rigorously:



# Markov Process, Revisited

We'll now define a Markov Process more rigorously:

## Markov Property

A Markov Process is a tuple  $\langle \mathcal{S}, \mathcal{P} \rangle$ :

- $\mathcal{S}$  is a finite set of states
- $\mathcal{P}$  is a state transition probability matrix


# Markov Process, Revisited

We'll now define a Markov Process more rigorously:

## Markov Property

A Markov Process is a tuple  $\langle \mathcal{S}, \mathcal{P} \rangle$ :

- $\mathcal{S}$  is a finite set of states
- $\mathcal{P}$  is a state transition probability matrix



# Reinforcement Learning Framework

Agent/Environment interaction in Reinforcement Learning is usually modeled as a *Markov Decision Process* (MDP from now on)

# Reinforcement Learning Framework

Agent/Environment interaction in Reinforcement Learning is usually modeled as a *Markov Decision Process* (MDP from now on)

- $s \in \mathcal{S}$  - state of the system

# Reinforcement Learning Framework

Agent/Environment interaction in Reinforcement Learning is usually modeled as a *Markov Decision Process* (MDP from now on)

- $s \in \mathcal{S}$  - state of the system
- $a \in \mathcal{A}$  - agent's action

# Reinforcement Learning Framework

Agent/Environment interaction in Reinforcement Learning is usually modeled as a *Markov Decision Process* (MDP from now on)

- $s \in \mathcal{S}$  - state of the system
- $a \in \mathcal{A}$  - agent's action
- $\mathcal{P} = p(s'|s, a)$  - the dynamics of the system

# Reinforcement Learning Framework

Agent/Environment interaction in Reinforcement Learning is usually modeled as a *Markov Decision Process* (MDP from now on)

- $s \in \mathcal{S}$  - state of the system
- $a \in \mathcal{A}$  - agent's action
- $\mathcal{P} = p(s'|s, a)$  - the dynamics of the system
- $\mathcal{R} : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$  - the reward function (possibly stochastic)

# Reinforcement Learning Framework

Agent/Environment interaction in Reinforcement Learning is usually modeled as a *Markov Decision Process* (MDP from now on)

- $s \in \mathcal{S}$  - state of the system
- $a \in \mathcal{A}$  - agent's action
- $\mathcal{P} = p(s'|s, a)$  - the dynamics of the system
- $\mathcal{R} : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$  - the reward function (possibly stochastic)

## Definition

**MDP** is the tuple  $\langle \mathcal{S}, \mathcal{P}, \mathcal{A}, \mathcal{R} \rangle$



# Reinforcement Learning Framework

Agent/Environment interaction in Reinforcement Learning is usually modeled as a *Markov Decision Process* (MDP from now on)

- $s \in \mathcal{S}$  - state of the system
- $a \in \mathcal{A}$  - agent's action
- $\mathcal{P} = p(s'|s, a)$  - the dynamics of the system
- $\mathcal{R} : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$  - the reward function (possibly stochastic)

## Definition

**MDP** is the tuple  $\langle \mathcal{S}, \mathcal{P}, \mathcal{A}, \mathcal{R} \rangle$

- $\pi(a|s)$  - the agent's policy
- A discount function/factor  $\gamma$  also usually exists

# Why Discount?

- Avoids infinite returns in cyclic Markov Chains
- In a lot of real life scenarios, immediate rewards may be of more interest than delayed awards
- ...as is common human/animal behavior

# Reinforcement Learning Framework

- It's common to break agent interaction into *episodes*
- Each episode begins with state  $s_0$ , drawn from some distribution  $\mu(s_0)$  and ends at some terminal state
- An action  $a_t$  will be chosen by the agent with accordance to the policy  $\pi$
- The next state is sampled according to  $\mathcal{P}$  presented earlier;  
 $\mathbb{P}(s_{t+1} \mid s_t, a_t)$
- Taking rewards into account, this becomes:

$$\mathbb{P}(s_{t+1}, r_t \mid s_t, a_t)$$

- The process continues until a terminal state is reached

# Trajectories and Episodes

- It's also common to introduce *trajectories*
- A trajectory  $\tau$  is simply a series of states ending at some timestep  $T$

$$\tau = \{(a_t, s_t)\}_{t=1}^T$$

- A trajectory can potentially contain multiple episodes

# Trajectories and Episodes

- It's also common to introduce *trajectories*
- A trajectory  $\tau$  is simply a series of states ending at some timestep  $T$

$$\tau = \{(a_t, s_t)\}_{t=1}^T$$

- A trajectory can potentially contain multiple episodes
- Those terms get confused a lot in the literature (and for iterative tasks are often interchangeable)

# Outline

## 1 What is Reinforcement Learning About

- Reinforcement Learning Examples
- Formal Introduction

## 2 Approaches to Reinforcement Learning

- Model-based RL
- Markov Decision Processes
- **Value-based RL**
- Bellman Equations
- Partially Observable Markov Decision Processes
- Policy-based RL

## 3 Solving RL in Practice

- Q-Learning
- Policy Gradient Methods
- Policy Gradients Explained
- REINFORCE
- Algorithm Comparison

## 4 Using Policy Gradients with Neural Networks

- Quick Example

## 5 Wrapping Up

- Further Reading/Acknowledgements

# The Value Function

- A **value function** is a prediction of future award
- Sometimes called *state-value function* - we ask "how much reward will I get from state  $s_t$  onwards?"
- The un-discounted value function for a policy  $\pi$ :

$$V^{\pi}(s) = \mathbb{E} \left[ \sum_{t=0}^T r(s_t, a_t) \mid s_0 = s \right]$$

- Where  $T$  denotes the *horizon*

# The Value Function

- A **value function** is a prediction of future award
- Sometimes called *state-value function* - we ask "how much reward will I get from state  $s_t$  onwards?"
- The un-discounted value function for a policy  $\pi$ :

$$V^\pi(s) = \mathbb{E} \left[ \sum_{t=0}^T r(s_t, a_t) \mid s_0 = s \right]$$

- Where  $T$  denotes the *horizon*
- The discounted value function is then:

$$V^\pi(s) = \mathbb{E} \left[ \sum_{t=0}^T \gamma(t) r(s_t, a_t) \mid s_0 = s \right]$$

- Where typically  $\gamma$  is simply a scalar  $0 < \gamma < 1$  so that  $\gamma(t) = \gamma^t$



# The Value Function - Goal

- Our goal is to maximize our value function; performed by finding the policy  $\pi^*$  which maximizes it  $\forall s \in \mathcal{S}$ :

$$V^*(s) = V^{\pi^*} = \max_{\pi} V^{\pi}(s)$$

# Outline

## 1 What is Reinforcement Learning About

- Reinforcement Learning Examples
- Formal Introduction

## 2 Approaches to Reinforcement Learning

- Model-based RL
- Markov Decision Processes
- Value-based RL
- **Bellman Equations**
- Partially Observable Markov Decision Processes
- Policy-based RL

## 3 Solving RL in Practice

- Q-Learning
- Policy Gradient Methods
- Policy Gradients Explained
- REINFORCE
- Algorithm Comparison

## 4 Using Policy Gradients with Neural Networks

- Quick Example

## 5 Wrapping Up

- Further Reading/Acknowledgements

# The Bellman Equation

- The value function can be written recursively:

$$\begin{aligned} V^\pi(s) &= \mathbb{E} \left[ \sum_{t=0}^T \gamma^t r(s_t, a_t) \mid s_0 = s \right] \\ &= \mathbb{E}_{\substack{a \sim \pi(\cdot|s) \\ s' \sim p(\cdot|s,a)}} \left[ r(s, a) + \gamma V^\pi(s') \right] \end{aligned}$$

- The optimal value satisfies the Bellman equation:

$$V^*(s) = \max_{\pi} \mathbb{E}_{\substack{a \sim \pi(\cdot|s) \\ s' \sim p(\cdot|s,a)}} \left[ r(s, a) + \gamma V^*(s') \right]$$

- Proof exists in literature

# Q-Function

- The  $Q$ -function gives the expected total reward of
  - from state  $s_t$  and action  $a_t$
  - under policy  $\pi$
  - with discount factor  $\gamma$

$$Q(s, a) = r(s, a) + \gamma \mathbb{E}_{s' \sim p(\cdot, s, a)} [V(s')]$$

# Q-Function

- The  $Q$ -function gives the expected total reward of
  - from state  $s_t$  and action  $a_t$
  - under policy  $\pi$
  - with discount factor  $\gamma$

$$Q(s, a) = r(s, a) + \gamma \mathbb{E}_{s' \sim p(\cdot, s, a)} [V(s')]$$

- If we know  $V^*$ , the optimal policy is to deterministically progress:

$$\pi^*(s) = \operatorname{argmax}_a Q^*(s, a)$$

- This implies a very important relationship between  $V$  and  $Q$ :

$$V^*(s) = \max_a Q^*(s, a)$$

# Optimal Value Functions

- An optimal value function is the maximum achievable value

$$Q^*(s, a) = \max_{\pi} Q^{\pi}(s, a) = Q^{\pi^*}(s, a)$$

- ...at *all* states:

$$\begin{aligned} Q^*(s, a) &= r_{t+1} + \gamma \max_{a_{t+1}} r_{t+2} + \gamma^2 \max_{a_{t+2}} + \dots \\ &= r_{t+1} + \gamma \max_{a_{t+1}} Q^*(s_{t+1}, a_{t+1}) \end{aligned}$$

- Again, we get the Bellman equation:

$$Q^*(s, a) = \mathbb{E}_{s'} \left[ r + \gamma \max_{a'} Q^*(s', a') \mid s, a \right]$$

# Outline

## 1 What is Reinforcement Learning About

- Reinforcement Learning Examples
- Formal Introduction

## 2 Approaches to Reinforcement Learning

- Model-based RL
- Markov Decision Processes
- Value-based RL
- Bellman Equations
- **Partially Observable Markov Decision Processes**
- Policy-based RL

## 3 Solving RL in Practice

- Q-Learning
- Policy Gradient Methods
- Policy Gradients Explained
- REINFORCE
- Algorithm Comparison

## 4 Using Policy Gradients with Neural Networks

- Quick Example

## 5 Wrapping Up

- Further Reading/Acknowledgements

# Fully Observeable Setting

- The framework described above pertain to the *fully-observable setting* - where the agent can observe the full dynamics of the system
- A more powerful framework exists, in which the agent only has access to an observation  $o_t$  at a given timestep



# Fully Observable Setting

- The framework described above pertain to the *fully-observable setting* - where the agent can observe the full dynamics of the system
- A more powerful framework exists, in which the agent only has access to an observation  $o_t$  at a given timestep
- In this instance, the true state  $s_t$  becomes a summary of experience:

$$s_t = f(o_1, r_1, a_1, \dots, a_{t-1}, o_t, r_t)$$

- If  $f(o_t) = s_t$ , this reduces back to the observable setting
- The more powerful framework is called a *Partially Observable Markov Decision Process* (POMDP)

# POMDPs

- Solving POMDPs hard!

# POMDPs

- Solving POMDPs hard!
- POMDP can be viewed as a natural extension of HMMs similarly to defining MDPs though Markov Chains
- One way of reducing a POMDP to an MDP is by defining a timeframe of last  $k$  steps as the agents current state - denoted *history* in the literature
- This results in:

$$a_{t+1} \sim \pi(a_t \mid \mathbf{h}_t)$$

- One common way of modeling  $\mathbf{h}_t$  in practice is to use an RNN
- We'll stick to MDPs today

# Outline

## 1 What is Reinforcement Learning About

- Reinforcement Learning Examples
- Formal Introduction

## 2 Approaches to Reinforcement Learning

- Model-based RL
- Markov Decision Processes
- Value-based RL
- Bellman Equations
- Partially Observable Markov Decision Processes
- Policy-based RL

## 3 Solving RL in Practice

- Q-Learning
- Policy Gradient Methods
- Policy Gradients Explained
- REINFORCE
- Algorithm Comparison

## 4 Using Policy Gradients with Neural Networks

- Quick Example

## 5 Wrapping Up

- Further Reading/Acknowledgements

# Policy

- The **Policy** is the agent's behavior
- Usually represented by a mapping from current state to action -  $\pi$ 
  - Deterministic:  $a_t = \pi(s_t)$
  - ...or Stochastic:  $\pi(a | s) = \mathbb{P}(a | s)$
- Policy-based RL is about searching directly for an optimal policy  $\pi^*$
- This is the policy which achieves the maximum future reward

# Outline

## 1 What is Reinforcement Learning About

- Reinforcement Learning Examples
- Formal Introduction

## 2 Approaches to Reinforcement Learning

- Model-based RL
- Markov Decision Processes
- Value-based RL
- Bellman Equations
- Partially Observable Markov Decision Processes
- Policy-based RL

## 3 Solving RL in Practice

- Q-Learning
- Policy Gradient Methods
- Policy Gradients Explained
- REINFORCE
- Algorithm Comparison

## 4 Using Policy Gradients with Neural Networks

- Quick Example

## 5 Wrapping Up

- Further Reading/Acknowledgements

# RL Methods by Type

RL is exploding in current research, but most methods used are pretty old. We will focus on two popular methods;

- $Q$ -Learning - learn a value function which satisfies the Bellman equations
- Policy Gradient - attempt to directly learn a policy maximizing rewards (using an analytic iterative solution, hence the "gradient")

# RL Methods by Type

RL is exploding in current research, but most methods used are pretty old. We will focus on two popular methods;

- $Q$ -Learning - learn a value function which satisfies the Bellman equations
- Policy Gradient - attempt to directly learn a policy maximizing rewards (using an analytic iterative solution, hence the "gradient")
- These are all *model-free* methods



# Model-Free Methods

- For the Atari example of DQN, states are  $64 \times 64$  RGB images and there are 4 actions

# Model-Free Methods

- For the Atari example of DQN, states are  $64 \times 64$  RGB images and there are 4 actions
- The transition dynamics  $\mathcal{P}$  are of size:

$$|S \times S \times A| \approx (4.68 \times 10^{1954}) \times (4.68 \times 10^{1954}) \times 4$$

# Model-Free Methods

- For the Atari example of DQN, states are  $64 \times 64$  RGB images and there are 4 actions
- The transition dynamics  $\mathcal{P}$  are of size:

$$|S \times S \times A| \approx (4.68 \times 10^{1954}) \times (4.68 \times 10^{1954}) \times 4$$

- Model-free implies we make no effort to learn the underlying dynamics of the environment
- Instead, we estimate the policy/value-function directly by interacting with the environment

# RL Methods by Type

Many other methods exist:

- Value Based- try to learn a value function satisfying Bellman Equation
  - *Q-Learning*, Double-*Q*-Learning
  - Temporal Difference (TD) Learning
  - SARSA
- Policy Search- attempt to directly learn a policy maximizing rewards
  - Gradient Based
    - *Policy Gradient methods* (Natural Policy Gradients, REINFORCE)
    - Actor-Critic algorithms (AC3)
  - Gradient-Free
    - Simulated Annealing
    - Cross-Entropy Search

# Outline

## 1 What is Reinforcement Learning About

- Reinforcement Learning Examples
- Formal Introduction

## 2 Approaches to Reinforcement Learning

- Model-based RL
- Markov Decision Processes
- Value-based RL
- Bellman Equations
- Partially Observable Markov Decision Processes
- Policy-based RL

## 3 Solving RL in Practice

- Q-Learning**
- Policy Gradient Methods
- Policy Gradients Explained
- REINFORCE
- Algorithm Comparison

## 4 Using Policy Gradients with Neural Networks

- Quick Example

## 5 Wrapping Up

- Further Reading/Acknowledgements

# Learning Optimal $Q^*$

- If the agent knows the dynamics  $p$  and the reward function  $r$ , it can find  $Q^*$  by dynamic programming
  - Many methods exist!
  - ...but they are useless if we don't know the dynamics, or if the state-space is huge
- Otherwise, it needs to estimate  $Q^*$  from its experience
  - The *experience* of an agent is a sequence  $s_0, a_0, r_0, s_1, a_1, r_1, \dots$
  - The  $n$ -th time-step is  $(s_n, a_n, r_n, s_{n+1})$

# Simulation-Based $Q$ -Value Iteration

- Lets interact with the environment for  $T$  timesteps and obtain a trajectory  $(s_0, a_0, \dots, s_T)$  according to policy  $\pi^{(n-1)}$ 
  - $\pi^{(n-1)}$  can be based on  $Q_{n-1}$
- Now update  $Q_n(s, a)$  to better represent our sample

# Off-Policy vs. On-Policy

- The update can use some  $a'$  drawn from *some* policy
- I.e., it doesn't depend on what action we actually took at state  $s_{t+1}$
- This allows us to sample our environment w.r.t. any policy - not only the one we are learning (e.g.  $\pi^{(n)}$ )
- This is called an *off-policy* approach



# Moving to an Online Algorithm

- The general algorithm works for batches/entire trajectories
- What if we go with  $n = 1$ ?

# Moving to an Online Algorithm

- The general algorithm works for batches/entire trajectories
- What if we go with  $n = 1$ ?
- We'll receive an *online* algorithm instead!
- Lets take an estimate of the environment based only on current state and move forward
- This results in the *Q*-Learning algorithm [WatkinsDayan92]

# $Q^*$ -Learning Algorithm

The full  $Q$ -Learning algorithm:

## $Q$ -Learning

**input:** learning rate  $\alpha \in (0, 1)$

**initialize:**  $Q_0(s, a)$  for all  $a \in \mathcal{A}$  and  $s \in \mathcal{S}$

**for each** time-step  $n$  **do**

observe the current state  $s_n$

select and execute an action  $a_n$

receive reward  $r_n$

observe the next state  $s_{n+1}$

**let:**

$$Q_n(s_n, a_n) \leftarrow (1 - \alpha_n)Q_{n-1}(s_n, a_n) + \alpha_n \left( r_n + \gamma \max_a Q_{n-1}(s_{n+1}, a) \right)$$

**end for**

# Update Step Explained

Lets dissect the update step:

## Q-Learning

$$Q(s_n, a_n) \leftarrow \underbrace{(1 - \alpha_n)Q_{n-1}(s_n, a_n)}_{\text{old value}} + \underbrace{\alpha_n}_{\text{learning rate}} \cdot \left( \underbrace{r_n}_{\text{reward}} + \underbrace{\gamma}_{\text{discount factor}} \cdot \underbrace{\max_a Q_{n-1}(s_{n+1}, a)}_{\text{estimate of optimal future value}} \right)$$

# Why Not $V$ -Learning?

- One might ask why we don't see  $V$ -Learning in the literature
- $Q$ -values make actions explicit
- Thus, they work when the transition function is not available
- As is usually the case for most interesting problems

# Exploration vs. Exploitation

- How should the agent gain experience in order to learn?
  - If it explores too much it might only run into bad options
  - If it exploits learned  $Q$  function too early (or too frequently) it might not learn about better ones!
- This is the RL *Exploration vs. Exploitation* problem
- $\epsilon$ -greedy policy - try a new action regardless of current learned policy with probability  $\epsilon$

# Credit Assignment Problem

- How does one know which action contributed to a high award for a given episode?
  - An action can have an effect in the far future
- This is the RL *Credit Assignment* problem
- Different solutions exist; decaying reward function is one
- The hand-wavvy reasoning behind this is that repeating the same task hundreds of thousands of times eliminates the issue

# Q-Learning - Wrap Up

- Q-Learning algorithm has been proven to find an optimal result for finite MDPs [Watkins and Dayan, 92]
- Model free - only uses Q-function and not system dynamics  
 $\mathbb{P}(s_{t+1}|s_t, a_t)$
- Difficult to adapt to non-Markovian settings
  - Model as a POMDP instead?
  - Partial observability makes the learning problem much harder and often intractable
- Also problematic for continuous or very large action-state spaces
  - Function approximators assist there
- Fuses beautifully with NNs (what would be the loss?)



# Outline

## 1 What is Reinforcement Learning About

- Reinforcement Learning Examples
- Formal Introduction

## 2 Approaches to Reinforcement Learning

- Model-based RL
- Markov Decision Processes
- Value-based RL
- Bellman Equations
- Partially Observable Markov Decision Processes
- Policy-based RL

## 3 Solving RL in Practice

- Q-Learning
- **Policy Gradient Methods**
- Policy Gradients Explained
- REINFORCE
- Algorithm Comparison

## 4 Using Policy Gradients with Neural Networks

- Quick Example

## 5 Wrapping Up

- Further Reading/Acknowledgements

# Learning the Policy Directly

- Optimize a policy end-to-end by computing an estimate of the *gradient of the expected reward of the policy*
- Assume a stochastic policy  $\mu(a_t | s_t)$  - giving a prob. distribution over actions

# Learning the Policy Directly

- Optimize a policy end-to-end by computing an estimate of the *gradient of the expected reward of the policy*
- Assume a stochastic policy  $\mu(a_t | s_t)$  - giving a prob. distribution over actions
- Optimally, examples with high rewards for *good actions* and low rewards for *bad actions* would result in increasing probability of good action
- Vanilla PG runs into a lot of problems and is rarely used in practice
  - REINFORCE is a popular solution

# Learning the Policy Directly

- Broadly speaking, if our policy is chosen w.r.t to parameter-set  $\theta$ , we want  $\max_{\theta} \mathbb{E}[R \mid \theta]$ 
  - Where  $R$  is the total reward of an episode
  - ...and actions are chosen from  $\pi(a_t \mid s_t; \theta)$

# Learning the Policy Directly

- Broadly speaking, if our policy is chosen w.r.t to parameter-set  $\theta$ , we want  $\max_{\theta} \mathbb{E}[R \mid \theta]$ 
  - Where  $R$  is the total reward of an episode
  - ...and actions are chosen from  $\pi(a_t \mid s_t; \theta)$
- Find gradient of policy w.r.t. current parameter-set  $\theta$
- Optimize using SGD

# Credit Assignment Problem, Again

- We should all be familiar with MLE optimization of this form:

$$\max_{\theta} \sum_{n=1}^N \log p(y_n | x_n; \theta)$$

# Credit Assignment Problem, Again

- We should all be familiar with MLE optimization of this form:

$$\max_{\theta} \sum_{n=1}^N \log p(y_n | x_n; \theta)$$

- Where we maximize the log-probability of outputs  $y_n$  for inputs  $x_n$

# Credit Assignment Problem, Again

- We should all be familiar with MLE optimization of this form:

$$\max_{\theta} \sum_{n=1}^N \log p(y_n | x_n; \theta)$$

- Where we maximize the log-probability of outputs  $y_n$  for inputs  $x_n$
- Here we instead wish to optimize action  $a_t$  for states  $s_t$
- Had we known  $a^*$  - the optimal action for each state; we could simply optimize:

$$\max_{\theta} \sum_{n=1}^N \log p(a_n^* | s_n; \theta)$$



# Credit Assignment Problem, Again

- We should all be familiar with MLE optimization of this form:

$$\max_{\theta} \sum_{n=1}^N \log p(y_n | x_n; \theta)$$

- Where we maximize the log-probability of outputs  $y_n$  for inputs  $x_n$
- Here we instead wish to optimize action  $a_t$  for states  $s_t$
- Had we known  $a^*$  - the optimal action for each state; we could simply optimize:

$$\max_{\theta} \sum_{n=1}^N \log p(a_n^* | s_n; \theta)$$

- Alas, we don't

# Outline

## 1 What is Reinforcement Learning About

- Reinforcement Learning Examples
- Formal Introduction

## 2 Approaches to Reinforcement Learning

- Model-based RL
- Markov Decision Processes
- Value-based RL
- Bellman Equations
- Partially Observable Markov Decision Processes
- Policy-based RL

## 3 Solving RL in Practice

- Q-Learning
- Policy Gradient Methods
- **Policy Gradients Explained**
- REINFORCE
- Algorithm Comparison

## 4 Using Policy Gradients with Neural Networks

- Quick Example

## 5 Wrapping Up

- Further Reading/Acknowledgements

# Approximating Good vs Bad Actions

- For a given trajectory

$$\tau = (s_0, a_0, r_0, s_1, a_1, r_1, \dots, s_{T-1}, a_{T-1}, r_{T-1}, s_T)$$

- Make a 'guess' at which actions were good and which weren't
- Increase probability of good actions repeating

# Approximating Good vs Bad Actions

- For a given trajectory
$$\tau = (s_0, a_0, r_0, s_1, a_1, r_1, \dots, s_{T-1}, a_{T-1}, r_{T-1}, s_T)$$
- Make a 'guess' at which actions were good and which weren't
- Increase probability of good actions repeating
- Let  $R = \sum_{t=0}^{T-1} r_t$  the sum of rewards

# Approximating Good vs Bad Actions

- For a given trajectory  
 $\tau = (s_0, a_0, r_0, s_1, a_1, r_1, \dots, s_{T-1}, a_{T-1}, r_{T-1}, s_T)$
- Make a 'guess' at which actions were good and which weren't
- Increase probability of good actions repeating
- Let  $R = \sum_{t=0}^{T-1} r_t$  the sum of rewards
- Assuming we can get  $\nabla_{\theta} \mathbb{E}[R]$  we are done!

# Getting a Derivative of the Policy

- We need to get a derivative for the policy function w.r.t. to some weights  $\theta$
- We don't actually know which actions resulted in outcomes
- We will use some math trickery to get a gradient-descent step which updates actions w.r.t. their 'usefulness'
- Lets name this mythical gradient  $\hat{g}$

# Training in Practice

- Assume the policy is some parametric model whose parameter set is  $\theta$
- Interact with environment 100 times, with 200 steps for each such interaction
- In total, we have 20,000 environment interactions
- Update every interaction which resulted in a 'bad' outcome as 'bad' and vice-versa
- This doesn't make sense for a short period, but does for a long series of such interactions

# Formal Explanation

- Idea based on *score function gradient estimator*
- Used for expressions of the form  $\mathbb{E}_{x \sim p(x|\theta)}[f(x)]$



# Formal Explanation

- Idea based on *score function gradient estimator*
- Used for expressions of the form  $\mathbb{E}_{x \sim p(x|\theta)}[f(x)]$ 
  - Expectation of scalar valued *score function*  $f$
  - For  $x$  drawn from  $p$  w.r.t  $\theta$
  - In our case,  $f$  is our reward function (was  $R$  earlier)
  - ...and  $p$  is the policy  $\pi(a | s; \theta)$

# Formal Explanation

- Idea based on *score function gradient estimator*
- Used for expressions of the form  $\mathbb{E}_{x \sim p(x|\theta)}[f(x)]$ 
  - Expectation of scalar valued *score function*  $f$
  - For  $x$  drawn from  $p$  w.r.t  $\theta$
  - In our case,  $f$  is our reward function (was  $R$  earlier)
  - ...and  $p$  is the policy  $\pi(a | s; \theta)$
- Math!

# Likelihood Ratio Trick

## Likelihood Ratio Trick

$$\begin{aligned}
 \nabla_{\theta} \mathbb{E}_x[f(x)] &= \nabla_{\theta} \sum_x p(x; \theta) f(x) && \text{def. of expectation} \\
 &= \sum_x \nabla_{\theta} p(x; \theta) f(x) && \text{insert gradient} \\
 &= \sum_x p(x; \theta) \frac{\nabla_{\theta} p(x; \theta)}{p(x; \theta)} f(x) && \text{multiply and divide by } p(x; \theta) \\
 &= \sum_x p(x; \theta) \nabla_{\theta} \log p(x; \theta) f(x) && \text{note } \nabla_{\theta} \log(z) = \frac{1}{z} \nabla_{\theta} z \\
 &= \mathbb{E}_x[\nabla_{\theta} \log p(x; \theta) f(x)] && \text{def. of expectation}
 \end{aligned}$$

# Likelihood Ratio Trick

## Likelihood Ratio Trick

$$\begin{aligned}
 \nabla_{\theta} \mathbb{E}_x[f(x)] &= \nabla_{\theta} \sum_x p(x; \theta) f(x) && \text{def. of expectation} \\
 &= \sum_x \nabla_{\theta} p(x; \theta) f(x) && \text{insert gradient} \\
 &= \sum_x p(x; \theta) \frac{\nabla_{\theta} p(x; \theta)}{p(x; \theta)} f(x) && \text{multiply and divide by } p(x; \theta) \\
 &= \sum_x p(x; \theta) \nabla_{\theta} \log p(x; \theta) f(x) && \text{note } \nabla_{\theta} \log(z) = \frac{1}{z} \nabla_{\theta} z \\
 &= \mathbb{E}_x[\nabla_{\theta} \log p(x; \theta) f(x)] && \text{def. of expectation}
 \end{aligned}$$

$\nabla_{\theta} \log p(x; \theta)$  is called the *score function*.

# Formal Explanation, Cont.

- Using the equality  $\nabla_{\theta} \mathbb{E}_x[f(x)] = \mathbb{E}_x[\nabla_{\theta} \log p(x; \theta) f(x)]$
- We can sample values  $x \sim p(x; \theta)$  and compute the LHS (over  $N$  samples)
- Hence;  $\hat{g} = \frac{1}{N} \sum_{n=1}^N \nabla_{\theta} \log p(x_i; \theta) f(x_i)$
- Full derivation in the literature!

# Problems

- The estimator  $\hat{g}$  is generally very noisy
- Instead of increasing the probability of *good episodes* - we would like to increase the probability of *good actions*
- This is called *gradient bias* in the literature
- Various solutions exist
- PG usually converges to a local maxima (unlike *Q*-Learning's guaranteed global)

# Outline

## 1 What is Reinforcement Learning About

- Reinforcement Learning Examples
- Formal Introduction

## 2 Approaches to Reinforcement Learning

- Model-based RL
- Markov Decision Processes
- Value-based RL
- Bellman Equations
- Partially Observable Markov Decision Processes
- Policy-based RL

## 3 Solving RL in Practice

- Q-Learning
- Policy Gradient Methods
- Policy Gradients Explained
- **REINFORCE**
- Algorithm Comparison

## 4 Using Policy Gradients with Neural Networks

- Quick Example

## 5 Wrapping Up

- Further Reading/Acknowledgements

# REINFORCE

- REINFORCE  $\rightarrow$  REward Increment = Nonnegative Factor times Offset Reinforcement times Characteristic Eligibility
- It turns out that we can improve the above formula by lowering the variance of gradient-estimates



# REINFORCE

- REINFORCE  $\rightarrow$  REward Increment = Nonnegative Factor times Offset Reinforcement times Characteristic Eligibility
- It turns out that we can improve the above formula by lowering the variance of gradient-estimates
- Add stability to the training process
- In practice, REINFORCE often converges where vanilla PG won't

# REINFORCE Algorithm

## REINFORCE

```
initialize:  $\theta$  at random  
for each trajectory  $\tau = \{s_1, a_1, r_1, \dots, a_{T-1}, r_T\} \sim \pi_\theta$  do  
  for  $t = 1 \dots T - 1$  do  
    let:  $\theta \leftarrow \theta + \alpha \nabla_\theta \log \pi_\theta(s_t, a_t) \cdot v_t$   
  end for  
end for
```

Note that  $v_t$  represents a stabilized reward and substitutes  $R$  in the vanilla PG

# Why Is 'Stabilization' Important?

- A simple example would be to normalize (e.g. subtract mean and divide by variance) of the computed reward  $R_t$
- This would guarantee that the number of 'bad' and 'good' actions (in terms of gradient computations) would be equal

# Outline

## 1 What is Reinforcement Learning About

- Reinforcement Learning Examples
- Formal Introduction

## 2 Approaches to Reinforcement Learning

- Model-based RL
- Markov Decision Processes
- Value-based RL
- Bellman Equations
- Partially Observable Markov Decision Processes
- Policy-based RL

## 3 Solving RL in Practice

- Q-Learning
- Policy Gradient Methods
- Policy Gradients Explained
- REINFORCE
- **Algorithm Comparison**

## 4 Using Policy Gradients with Neural Networks

- Quick Example

## 5 Wrapping Up

- Further Reading/Acknowledgements

# Q-Learning vs. Policy Gradients

|  | Q-Learning   | REINFORCE                       |
|--|--|---------------------------------|
| Learning objective                       | Value function   | Policy parameters via gradient  |
| Policy Stochasticity                     | Convergence to essentially deterministic policy          | Explicitly stochastic           |
| Model                                    | Model free   | Model free                      |
| Markovity Assumption                     | MDP\POMDP  | Problem specific, not inherent  |
| Support for continuous\large state space | Possible via function approximation, difficult to train. | Yes                             |
| Main Practical Successes                 | DeepMind Atari playing net                               | Visual Attention, robot control |

- **Q-Learning better adapted to classic observable Markovian setting, PG-based learning more relevant the further problem is from that setting (hidden and continuous states).**

# Outline

## 1 What is Reinforcement Learning About

- Reinforcement Learning Examples
- Formal Introduction

## 2 Approaches to Reinforcement Learning

- Model-based RL
- Markov Decision Processes
- Value-based RL
- Bellman Equations
- Partially Observable Markov Decision Processes
- Policy-based RL

## 3 Solving RL in Practice

- Q-Learning
- Policy Gradient Methods
- Policy Gradients Explained
- REINFORCE
- Algorithm Comparison

## 4 Using Policy Gradients with Neural Networks

- Quick Example

## 5 Wrapping Up

- Further Reading/Acknowledgements

# Policy Gradients in Neural Networks

- Policy Gradients are best paired with function approximators - particularly Neural Networks
  - ...this is also true for  $Q$ -Learning (e.g.  $DQN$ )
- For example, consider an RNN in which output represents action probabilities
  - Softmax output  $a_t \in [k]$  for  $k$  actions
  - Hidden units of RNN  $\theta_t$  are policy parameterization.
  - Weights frozen for duration of episode and updated at the end
  - Naturally compatible with backpropagation-  $\nabla_{\theta} \log \pi_{\theta}(a_t | s_t)$  is the gradient of the corresponding RNN evaluated at timestep  $t$

# Outline

## 1 What is Reinforcement Learning About

- Reinforcement Learning Examples
- Formal Introduction

## 2 Approaches to Reinforcement Learning

- Model-based RL
- Markov Decision Processes
- Value-based RL
- Bellman Equations
- Partially Observable Markov Decision Processes
- Policy-based RL

## 3 Solving RL in Practice

- Q-Learning
- Policy Gradient Methods
- Policy Gradients Explained
- REINFORCE
- Algorithm Comparison

## 4 Using Policy Gradients with Neural Networks


- Quick Example

## 5 Wrapping Up

- Further Reading/Acknowledgements




# Recurrent Attention Model




- Original paper by (Mnih et al., 2015)
- Policy parameterized by RNN
- At each step 2 types of actions ( $l_t$  glimpse location and  $a_t$  classification) controlled by 2 sub-networks
- Goal is to learn stochastic policy  $\pi((l_t, a_t) | s_{1:t}; \theta)$  maximizing rewards

# Example: Recurrent Attention Model (Mnih et al., 2015)




- Trajectory given by  $s_{1:t} = x_1, l_1, a_1, \dots, x_{t-1}, l_{t-1}, a_{t-1}, x_t$
- At each step 2 types of actions ( $l_t$  glimpse location and  $a_t$  classification) controlled by 2 sub-networks
- Reward  $R = \sum_{t=1}^T r_t$  where  $r_T = 1$  for correct classification and 0 otherwise

# Example: Recurrent Attention Model (Mnih et al., 2015)



## Example: Recurrent Attention Model (Mnih et al., 2015) - Dynamic Environment



<http://www.cs.toronto.edu/~vmnih/docs/attention.mov>

- Same approach used to train agent to play simple game in dynamic environment

# Outline

## 1 What is Reinforcement Learning About

- Reinforcement Learning Examples
- Formal Introduction

## 2 Approaches to Reinforcement Learning

- Model-based RL
- Markov Decision Processes
- Value-based RL
- Bellman Equations
- Partially Observable Markov Decision Processes
- Policy-based RL

## 3 Solving RL in Practice

- Q-Learning
- Policy Gradient Methods
- Policy Gradients Explained
- REINFORCE
- Algorithm Comparison

## 4 Using Policy Gradients with Neural Networks

- Quick Example

## 5 Wrapping Up

- Further Reading/Acknowledgements

# Outline

## 1 What is Reinforcement Learning About

- Reinforcement Learning Examples
- Formal Introduction

## 2 Approaches to Reinforcement Learning

- Model-based RL
- Markov Decision Processes
- Value-based RL
- Bellman Equations
- Partially Observable Markov Decision Processes
- Policy-based RL

## 3 Solving RL in Practice

- Q-Learning
- Policy Gradient Methods
- Policy Gradients Explained
- REINFORCE
- Algorithm Comparison

## 4 Using Policy Gradients with Neural Networks

- Quick Example

## 5 Wrapping Up

- Further Reading/Acknowledgements

## Further Reading/Where I Stole This From

- Sutton, R. S., Mcallester, D., Singh, S., Mansour, Y. (1999). Policy Gradient Methods for Reinforcement Learning with Function Approximation.
- Karpathy, A. (2016). Deep Reinforcement Learning: Pong from Pixels (blogpost)
- Sutton, R. Barto, A. (1998). Reinforcement Learning: An Introduction (Book)
- Watkins, C.J.H., Dayan, P. (1992) Technical Note: Q-Learning
- OpenAI Gym RL Introduction
- Mnih. V et Al. (2014). Playing Atari with Deep Reinforcement Learning
- Grondman. I et Al. (2012). A Survey of Actor-Critic Reinforcement Learning: Standard and Natural Policy Gradients

## Further Reading/Where I Stole This From

...I also blatantly stole slides from:

- David Silver's RL Course from UCL
- David Silver's ICML2016 RL Workshop slides
- David Silver's NIPS2016 RL Workshop slides
- Generally anything by David Silver
- REINFORCE presentation by Ronen Tamary of HUJI
- $Q$ -Learning presentation by Noga Zaslavsky of HUJI
  - Both presented at HUJI's own DL seminar - good stuff there!



# Outline

## 1 What is Reinforcement Learning About

- Reinforcement Learning Examples
- Formal Introduction

## 2 Approaches to Reinforcement Learning

- Model-based RL
- Markov Decision Processes
- Value-based RL
- Bellman Equations
- Partially Observable Markov Decision Processes
- Policy-based RL

## 3 Solving RL in Practice

- Q-Learning
- Policy Gradient Methods
- Policy Gradients Explained
- REINFORCE
- Algorithm Comparison

## 4 Using Policy Gradients with Neural Networks

- Quick Example

## 5 Wrapping Up

- Further Reading/Acknowledgements

# Thank You